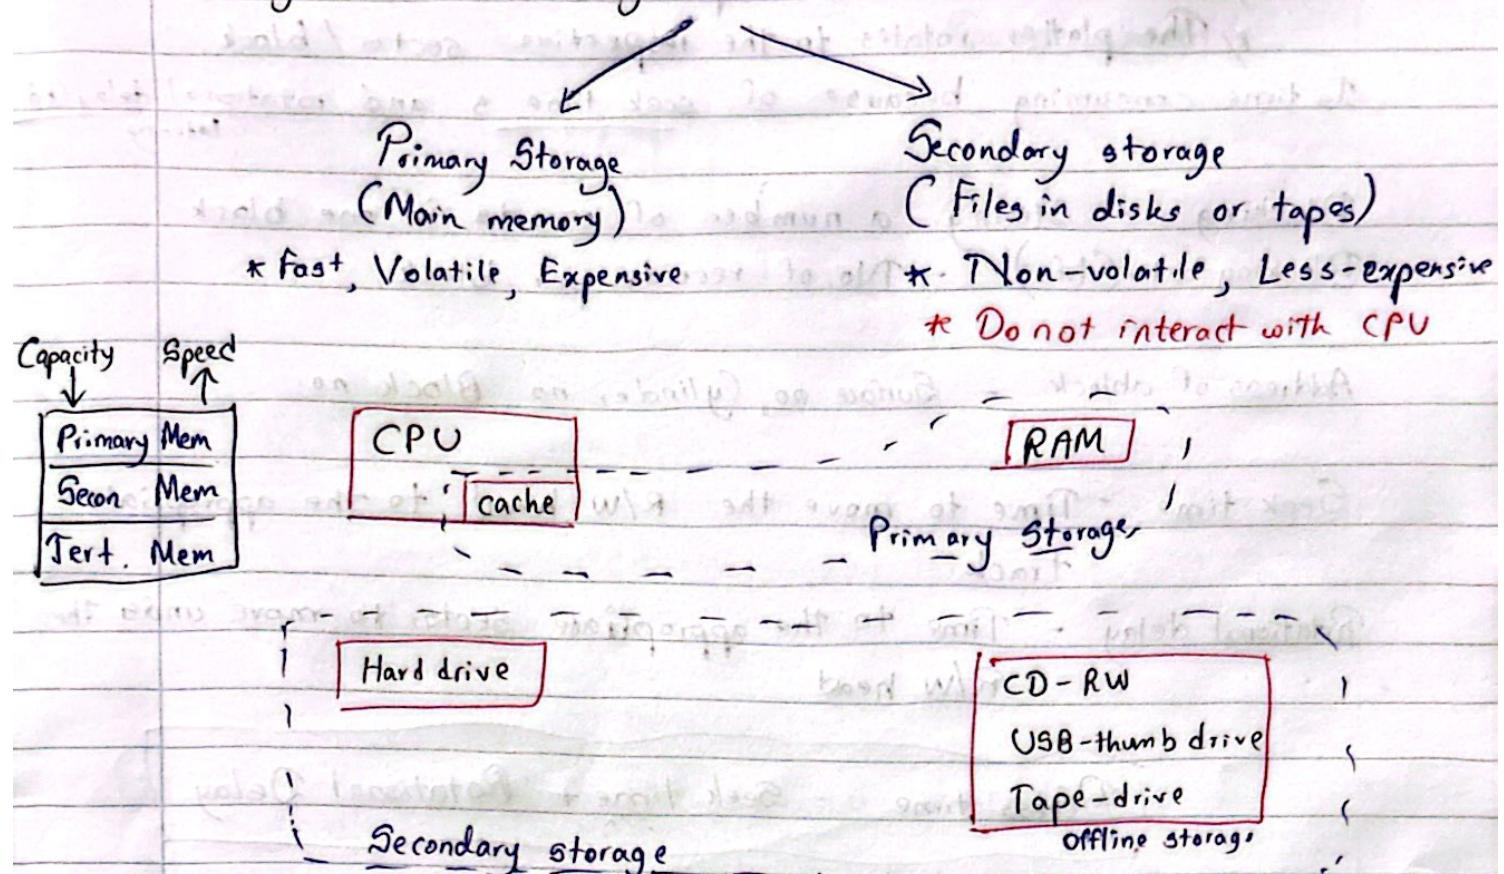


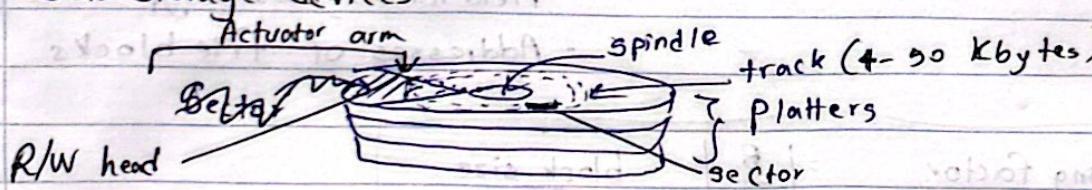
## Database Notes

### File organization and storage structures



- \* Databases are usually stored on magnetic disk secondary storage.
- \* Data stored in disks as files of records.
- \* Record is a collection of data values that can be interpreted as facts about entities, their attributes and their relationships.

### Disk storage devices



- \* Block size is fixed for each system.
- \* It is the smallest transferable unit.  
Vary from 512 bytes to 4096 bytes.

When transferring data,

i) R/W head moves to the respective track

ii) The platter rotates to the respective sector/block.

Is time consuming because of seek time & rotational delay (rd latency)

Blocking - Storing a number of records in one block

Blocking factor (bfm) - No. of records per block

Address of a block - Surface no, Cylinder no, Block no.

Seek time - Time to move the R/W head to the appropriate track

Rotational delay - Time to move the appropriate sector to move under the R/W head.

$$\boxed{\text{Access time} = \text{Seek time} + \text{Rotational Delay}}$$

Seek & Rotational time  $\ggg$  Transfer time

Bottleneck in

File - Sequence of records

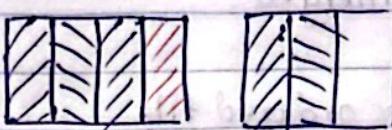
File descriptor/File header - Contains details about the file

- Field names, Data types
- Addresses of file blocks

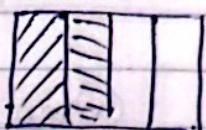
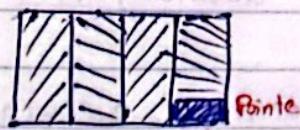
$$\boxed{\text{Blocking factor (bfm)} = \left\lfloor \frac{B}{R} \right\rfloor \cdot \frac{\text{block size}}{\text{record size}}}$$

## Record spanning

Unspanned records : Spanned records



free space



Spanned to the next sector

## Types of block allocation

### i, Contiguous allocation

- + Easy file scanning
- Difficult to expand the file by including blocks in middle

### ii, Linked allocation

- + Easy to expand
- Difficult to scan

### iii, Mixed allocation

## File operations - Read from slides

### Primary file organization methods

#### - Heap / Unordered file (Look at slides) Complexity ( $b/2$ )

\* No particular order

\* Appends new records at the end of the file

#### Normal - blz log b Binary search - - Ordered file (Sorted / Sequential) Complexity ( $\log b$ )

\* Records are ordered by a field value (sort key)

#### - Hashed file

+ Applies a hash function to a particular field to determine the placement

#### - B-trees

\* Use tree structures

When inserting a record in an ordered file

Insert the new record in an auxiliary file and when periodic file reorganization is happening, the main file and the auxiliary file will be merged.

When locating a record in an ordered file

First search is performed on the auxiliary file, if not found the main file is searched

### Linear ordering queries

- \* Create table sales

$$\text{prod\_id} \text{ int number(0) not null,}$$

$$\text{cust\_id} \text{ number not null}$$

$$)$$

CLUSTERING BY LINEAR ORDER (cust\_id, prod\_id);

- \* Alter table sales ADD CLUSTERING BY LINEAR ORDER(cust\_id, prod\_id)
- \* Alter table sales order by prod\_id ASC;

Heap / Unordered

Search Method:

Sequential

Average time to access:

b/2

Ordered

Sequential

b/2

Ordered

Binary

$\log_2 b$

Problem of a hash file

- \* Collisions occur when a new record hashes to a bucket that is already full. An overflow file is kept for storing such records.

## Indexing

To do these calculations, the records should be fixed size and

2) Slide 02 - Page 05

unspanned

$$bfr = \left\lceil \frac{1024}{100} \right\rceil = \lfloor 10.24 \rfloor = 10$$

$$\therefore \text{No. of blocks} = \frac{30000}{10} = \underline{\underline{3000}}$$

When optimizing query performance, the objective is to minimize the number of disk/block accesses. Because the performance depends on how many block accesses it takes to perform the query.

What is indexing?

Defining an additional data structure which provide an alternative path to access data without affecting its original representation to improve database operation's performance.

Indexing can be done as

- Creating an index using a single field.
- Creating multiple indexes on multiple fields.
- Indexes on multiple fields.

An index entry is of format

$\langle k(i), p(i) \rangle$

↗

↑

Search key

Block pointer / Row pointer

Indexing field - The field used to order the index records in the index file.

What could be the reason for the number of block accesses being less compared to data file?

- An index entry record length is comparably smaller than a data record.
- We can perform binary search on the index file always.

### Indexes

~~Sparse~~

\* Contains records for only some of the search key values.

~~Dense~~

\* Contains a record for every search key value.

Applicable to ordered data files only

### Single level indexes

Primary Index

Ordering key field of an ordered file

Clustering Index

Ordering non-key field of an ordered file

Secondary Index

Any non ordering field.

ordered file

May be key or non-key Field

\* ~~Ordered / Non ordered file.~~

## Primary Index (Ordered datafile, Ordered on a key field)

- \* Includes one index entry for each block in the datafile
- \* Index entry has the key field value for the first record in the block (block anchor)
- \* Sparse index.

$$\text{Accessing a record} = \log_2 b + 1$$

↑                      ↑  
Binary search      Accessing the block

Q) Slide 2, Page 21

$$bfr = \left\lceil \frac{1024}{100} \right\rceil = \lceil 10.24 \rceil = 10 \text{ records per block}$$

No. of block accesses to search

$$\text{No. of blocks} = \frac{3000}{10} = 300 \text{ blocks.}$$

$$\therefore \text{No. of block accesses to search} = \lceil \log_2 b \rceil$$

$$= \lceil \log_2 3000 \rceil$$

$$= \lceil 11.55 \rceil$$

$$= \underline{\underline{12}}$$

\*)

$$\text{Index record size} = 9 + 6 = 15 \text{ bytes}$$

$$\text{No. of records} : \text{No. of blocks} = 3000$$

$$bfr = \left\lceil \frac{1024}{15} \right\rceil = 68$$

$$\therefore \text{No. of block} = \lceil \log_2 45 \rceil$$

$$\text{No. of blocks} : \left\lceil \frac{3000}{68} \right\rceil = \frac{45}{68} \text{ block access} \approx 5.4$$

$$= \frac{6}{7}$$

Indexing Part II

Q) Slide 03, Page 04

$$\text{bfr} = \left\lceil \frac{2048}{80} \right\rceil = \left\lceil 25.6 \right\rceil = 26$$

$$\text{No. of blocks} : \frac{40000}{25} = 1600$$

Q) Page 06

\* )  $\text{bfr} = \left\lceil \frac{4096}{100} \right\rceil = 41$

$$\text{No. blocks} = \frac{300000}{40} = 7500$$

$$\text{Block accesses} = \log_2 7500 = 13$$

\* ) Index record size = 9 + 6 = 15 bytes

$$\text{bfr} = \left\lceil \frac{4096}{15} \right\rceil = 273$$

$$\text{No. of blocks} : \left\lceil \frac{7500}{273} \right\rceil = 28$$

$$\therefore \text{No. of block accesses for index} = \left\lceil \log_2 28 \right\rceil = 5$$

$$\text{No. of total block access} = 5 + 1 = 6$$

Problems of primary indexing

- \* Inserting and deleting records in the main file must move other records, since it's ordered

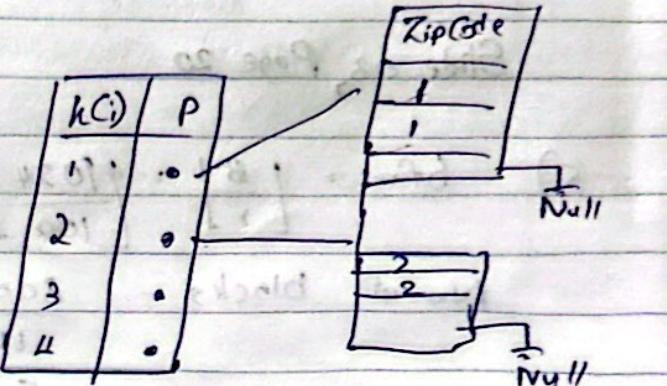
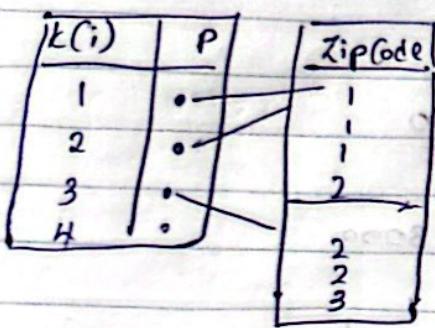
- \* Some insertions/deletions must also change index entries, the anchor records change.

- \* There can be only one primary index on a file.

## Clustering index (Ordered data file, Ordered non key field)

- \* Includes one index entry for each distinct value of the field
- \* Sparse index
- \* Index entry points to the first data block that contains record with that field value.

Two types of clustering index



\* New values start with the same block

\* New values start with ~~the~~<sup>new</sup> block

Page 15

Q) Index entry length =  $5 + 6 = 11$  bytes

$$bfr = \left\lceil \frac{4096}{11} \right\rceil = 372$$

$$\text{No. of blocks in indexfile} = \left\lceil \frac{1000}{372} \right\rceil = 3$$

$$\text{No. of block accesses} : \lceil \log_2 3 \rceil = 2$$

## Secondary index (Non ordering key/non key field)

\* There can be many secondary indexes defined for a single file.

If the index field is a key field (Secondary key)

↳ Then index is dense

If the index field is a non key field

↳ Then index is sparse

Slide 03, Page 20

$$\text{Q) } bfr = \left\lfloor \frac{B}{r} \right\rfloor = \left\lfloor \frac{1024}{100} \right\rfloor = 10$$

$$\text{No. of blocks} = \frac{30000}{100} = 3000$$

No. of block accesses =  $\lceil \log_2 3000 \rceil = 12 \times$  Cannot do binary as it is non ordered

$$\therefore \frac{3000}{2} = 1500$$

#) Index entry length =  $9+6 = 15$  bytes

~~Number of blocks = 30000~~

$$bfr = \left\lfloor \frac{1024}{15} \right\rfloor = 68$$

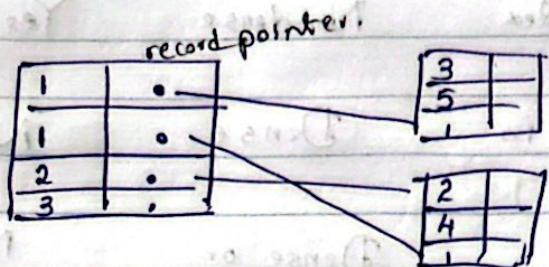
$$\text{No. of blocks} = \left\lceil \frac{30000}{68} \right\rceil = 442$$

$$\text{No. of block accesses} = \lceil \log_2 442 + 1 \rceil$$

$$= 10$$

If the field is a non-key field, which means there can be duplicate values, there are 3 ways that we can implement secondary index.

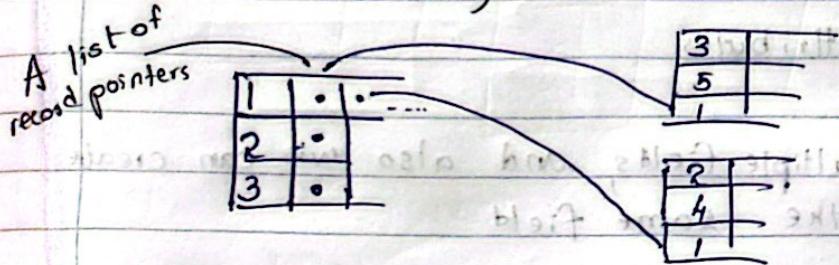
01/ Several indexes for same  $k(i)$  value, hence a dense index



\* In this case we have record pointers instead of block pointers.

\* Dense index

02/ One index for one  $k(i)$  value, but variable index record length

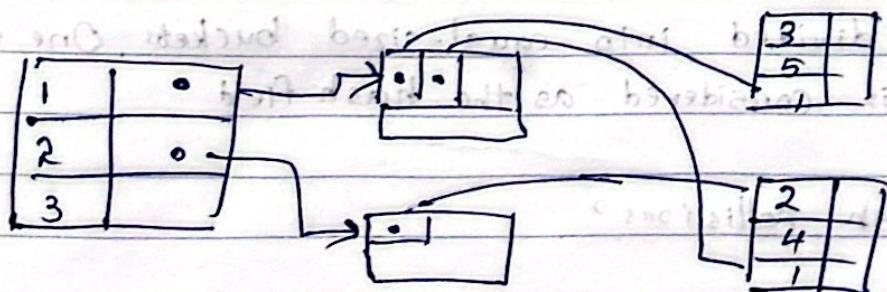


\* Record pointers and index records are

\* Variable length

\* Sparse index

03/ Have two indexes, the first index have fixed size length and a block pointer. The second index have a list of record pointers



\* Here the first index is of fixed length and it has block pointer

\* Non dense index

Type of Index	No. of first level index entries	Dense or Nondense	Block anchoring on the datafile
Primary Clustering	No. of blocks in data file	Nondense	Yes
Secondary (key)	No. of distinct index field values	Nondense	Yes / No
Secondary (nonkey)	No. of records in data file No. of distinct index values	Dense Nondense	No

### Indexing on multiple attributes

An index can have multiple fields, and also we can create multiple indexes for the same field.

Create index person\_fn\_idx on Person(first\_name, last\_name)

### Hash index

Blocks are divided into equal-sized buckets. One of the record field is considered as the hash field.

What is hash collisions?

When a new record is hashed to a bucket which is already full, it is known as a hash collision.

What are the hash collision resolution mechanisms?

- i) Open addressing - If the bucket given by hashing function is already full, move to the adjacent buckets until we find an empty slot.

$$\begin{aligned} h(18) \\ = 18 \bmod 5 \\ = 3^{\text{rd}} \text{ bucket} \end{aligned}$$

full	<table border="1"> <tr> <td>3</td><td>8</td><td>13</td></tr> </table>	3	8	13	3 <sup>rd</sup> bucket
3	8	13			
full	<table border="1"> <tr> <td>4</td><td>9</td><td>14</td></tr> </table>	4	9	14	4 <sup>th</sup> bucket
4	9	14			
full	<table border="1"> <tr> <td>5</td><td>18</td><td>10</td></tr> </table>	5	18	10	5 <sup>th</sup> bucket
5	18	10			

- ii) Chaining - Put the value in an overflow location and set a pointer to the intended location.

$$\begin{aligned} h(18) \\ = 18 \bmod 5 \\ = 3^{\text{rd}} \text{ bucket} \end{aligned}$$

fill → 

3	8	13
---	---	----

 → 

18
----

- iii) Multiple hashing - Apply a second hashing function if the first results in a collision.

\* Hash index is also a type of secondary index.

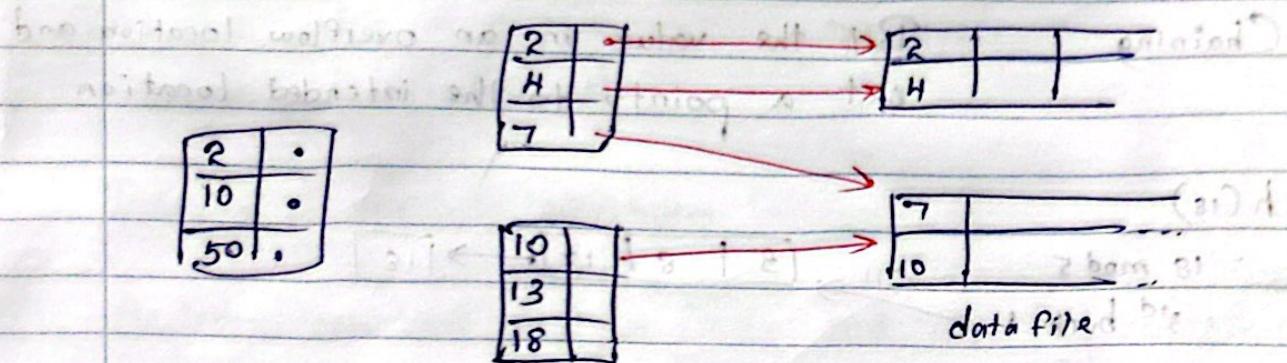
## Multi-level indexing

Since an index is ordered, we can create a primary index to the index file.

The initial index file is called the first-level index.

Index for the index is called the second-level index.

\* We can proceed with this approach until the ~~first~~ top level index fit into a single block.



Second index      First index

$$\text{Number of levels needed} = \lceil \log_{f_0} (r_1) \rceil$$

$f_0$  = fan out (blocking factor for index)

$r_1$  = Number of index entries in the first level index

A file organization that uses multi-level indexing is called indexed sequential file.

How do we handle insertion in multilevel indexing?

Create an overflow file, store it there and merge it to the main file periodically.

The indexes are recreated when reorganizing.

$$\textcircled{a}) \quad Bfr_1 = 68 \quad r_1 = 442 \times 68$$

$$\# \text{ of second level blocks} = \left\lceil \frac{442}{68} \right\rceil = 7$$

$$\# \text{ of third level blocks} = \left\lceil \frac{7}{68} \right\rceil = \underline{\underline{1}}$$

$\therefore$  There are 3 levels

$$\text{No. of } \underline{\underline{2}} \text{ block accesses} = 3 + 1 \\ = \underline{\underline{4}}$$